

Section 1: Contact information for the Israeli PI

Prof. Elisha Moses, Weizmann Institute, Phone: +917-8-934-3139, Email: elisha.moses@weizmann.ac.il

Section 2: IOS Program Justification

BIO/IOS Neural Systems NSF 11/572. Our proposed work is most appropriate for this section as it deals with fundamental questions relating to the computational capabilities of the central nervous systems of living organisms. The proposal links between theoretical and experimental methods of inquiry and addresses these in the context of co-evolution between the underlying biological organization and the computational tasks such structure is capable of performing in various sets of environmental conditions.

Section 3: Research description

Motivating questions

One of the big mysteries regarding the brain is, why neurons that make amazingly complex calculations in the brain become very inefficient at computation once they are grown out of it. In particular, once you take hippocampal neurons out of the rat brain and culture them on a dish, they lose their individuality. Living neuronal cultures are characterized by all-or-none network bursts, where practically all the neurons fire simultaneously, with varying degrees of synchrony. While in the brain they can compute the trajectory of the animal in a maze, but as an ensemble grown out of it they are unable to create new computations, and can only carry very few bits of information. How is it then that individual neurons are organized in the brain as a splendid computer, but in the dish their computational repertoire becomes very limited. Obviously, their environment and growth conditions have changed, but can we isolate and pin down those elements that are necessary, and perhaps also sufficient, to support computation by a living neuronal network? In an effort to tackle such questions, we have embarked on an investigation of the computational abilities of living neuronal networks, and propose here to expand this investigation both theoretically and experimentally in several fundamental ways.

On the conceptual side, a number of deep questions arise. What are the structural properties at the level of networks of neurons, modules of networks of neurons, and perhaps higher order forms of organization necessary to support the capacity for **abstraction**, which is fundamental to computation and may likewise be fundamental to **cognitive architecture**, development, and function [1]? Given that boolean logic devices are capable of being implemented in vitro using neurons [2], is it possible to biologically engineer analogous devices capable of performing computations that have natural embeddings within **first** or **higher order logic**? Moreover, can we define environmental conditions that lead to the development of an extensionally equivalent machine, but whose architecture may differ from the human-engineered biological implementation? Would multiple developmental implementations conserve identifiable structural features? And, how do the structural properties of networks with the capacity for zeroth, first, or higher order computations compare?

On the experimental side, we plan to design a number of novel logical devices, as well as use a number of new technologies that we have acquired for the construction of such computational constructs. Notable among these are the optogenetic toolbox that has recently become available. Light induced neuronal excitation using channelrhodopsins (courtesy of the Deisseroth lab in Stanford and the Yizhar lab at Weizmann) as well as genetically encoded fluorescent labels for neuronal excitation (courtesy of the Cohen lab at Harvard).

Finally, it seems clear that evolutionary forces have shaped the architecture, connectivity as well as

the input that the neurons grow with inside the brain, all of which conspire to create its tremendous computational power. Our most novel proposal is that biological computation in general is set apart from that of electronic computers by the fact that the 'hardware' (e.g. the cell or the organism) co-evolved with the 'software' (DNA or the brain respectively). As a result, our 'software' naturally tends to the well being of the associated 'hardware', while in the computer world that is not usually the case. This insight leads us to propose experiments that will connect the neuronal network with the 'real' world, and allow it to co-evolve with a range of different neurons and of hardware.

Physical implementations of Turing machines and other abstractly characterized computation

Since the development of the electronic digital computer, which makes use of the digital abstraction [3] from analog electrical circuits there has been a close heuristic association between boolean logic and computation. However, developments in formal logic over the past century have been largely motivated by its applications to computation and the theory of programming languages that sometimes build upon and sometimes go beyond boolean logic to support additional forms of abstraction. In fact, the capacity to support abstraction mirroring various systems of formal logic has become a means to order and thereby assign value to programming languages [4]. Electronic devices that are capable of supporting such forms of abstraction provide a concrete physical instantiation of the ideas inherent to the logical systems they are designed to faithfully implement.

Neuronal logic devices (NLDs) represent an alternative physical modality to digital electronics for the purpose of performing computation. However, rather than attempting to directly parallel the history of the development of digital electronics via the digital abstraction, high-level descriptions of computation, such as the λ -calculus, serve as a specification of computation that is agnostic to the physical modality of implementation. If the specification that a neuronal computation device should be capable of implementing the λ -calculus, a natural first step toward this broad goal is to investigate simple neuronal systems that are capable of performing well-defined computations that require a capacity for first- or higher-order logic.

The first question to be addressed experimentally is whether a universal Turing machine can reliably be built out of central nervous system (CNS) neurons, which are essentially unreliable. We have shown previously that complex devices such as a diode, an oscillator and an AND-gate can be engineered using CNS neurons grown on particular geometric configurations. The production of a NOT-gate would complement this set and enable the construction of a universal Turing machine. The design we are proposing initially relies on the fact that NLD's can be strung together to form more complex structures. To form a NOT-gate, a delay line, oscillator AND gate and diode can be put together in a way that ensures that the output is continually excited (at the frequency of the oscillator) unless the Input is high. This relies on the fact that living neuronal networks have a latency period following their spike. The NOT gate can then be extended to build a NAND gate, where if excitation comes from either one of the two Input regions then it can propagate and excite the Output region, but if both fire together then one cancels the other.

A higher-order function to be implemented as a higher-order NLD

To design higher order computational devices, we will turn to the experimental techniques of microfluidics and of optogenetics. It is now possible to monitor with optical means the electrical activity of the network without any collateral damage caused by the fluorescent dye. This is done by incorporating genetically a fluorescent voltage-indicating protein into the neuron. It is furthermore possible to excite a region of the network optically by activating photosensitive channels that are genetically embedded as well. On top of this, the propagation velocity of a signal inside a one-

dimensional neuronal network of the type we are using is constant, and can be reliably predicted. Thus it becomes possible to identify activity in one part of the devices, and then excite another region co-incidentally with the arrival of the signal into that area. Different time delays, with the activation before, during or after the arrival of the synaptic input will allow the creation of several neuronal learning scenarios, and the comparison to learning in the animal itself.

On the part of the theory, the λ -calculus notation is helpful in order to define any higher order function. Some of the notation may be implicitly familiar to users of imperative programming languages under the heading “anonymous functions”. We provide only an extremely informal set of examples necessary to explain our intended work; however, complete details can be found in [5]. We can define a standard binary boolean function like the “and” function as

$$\lambda x.\lambda y.(\wedge x y) \quad (1)$$

Such a lambda expression can apparently be applied to any inputs; however the inputs to such a function are not necessarily restricted to booleans unless we infer that the standard logical operator “ \wedge ” only accepts boolean arguments. Note that we have used the **prefix or Polish notation** for the \wedge operator, which in the perhaps more common infix notation is written with its arguments flanking the operator as $x \wedge y$ to mean “ x and y ”. We can provide explicit type annotations for the bound variables x and y , which indicate the types of the arguments

$$\lambda x : \text{bool}.\lambda y : \text{bool} . (\wedge x y) \quad (2)$$

We can now also provide a type annotation for this expression as a whole

$$[\lambda x : \text{bool}.\lambda y : \text{bool} . (\wedge x y)] : [\text{bool} \rightarrow \text{bool} \rightarrow \text{bool}] \quad (3)$$

Depending upon conventions with respect to **currying**, one could read the type annotation as “the function that takes two arguments each of type *bool* and returns a values of type *bool*” or “the function that takes an argument of type *bool* and returns a function that takes an argument of type *bool* and returns a value of type *bool*”. The first formulation may be easier to read, but the second is standard as a result of the way in which functional programming languages implement such functions.

Now we can imagine that if we wish to abstract from the particular binary boolean function implied by the \wedge operator, we need to introduce a functional variable, whose type will be explicitly denoted for concreteness despite the fact that it could be inferred, for which this operator can be substituted. Doing this results in the following second order function

$$\begin{aligned} [\lambda f : (\text{bool} \rightarrow \text{bool} \rightarrow \text{bool}).\lambda x : \text{bool}.\lambda y : \text{bool} . (f x y)] \\ : [(\text{bool} \rightarrow \text{bool} \rightarrow \text{bool}) \rightarrow \text{bool} \rightarrow \text{bool} \rightarrow \text{bool}] \end{aligned} \quad (4)$$

This function is intended to be read, in the less verbose uncurried form, as “the function that takes as its first argument (a function that takes two boolean values as arguments and returns a boolean value) and as its second and third arguments two boolean values and returns a boolean value”. It is perhaps remarkable that this simple abstraction is now capable of implementing any of the 16 possible boolean functions, provided that the proper binary boolean operator is submitted as the first argument to this function. Furthermore, the statement of this function in terms of a simply typed lambda calculus notation is agnostic to any physical implementation capable of realizing extensionally equivalent behavior.

To make this more concrete, we can very simply implement the above function in any functional programming language, or any programming language supporting the passing of function handles to other functions. An implementation in the programming language OCaml appears as follows

```
1 let hobf = fun (bf : ('a 'a bool)) (i1 : bool) (i2 : bool)
    -> bf i1 i2
```

Listing 1: a higher order boolean function

What is required in order to evaluate this function are corresponding implementations of boolean functions to be substituted either for f in the lambda calculus notation or for bf in terms of the corresponding OCaml implementation. For example we can implement the XOR function using pattern matching to define a truth table.

```
let xor p1 p2 = match (p1, p2)
2 with (false, false) -> false
    | (false, true) -> true
4   | (true, false) -> true
    | (true, true) -> false
```

Listing 2: implementation of an XOR boolean operator

Other binary boolean functions are implemented in an analogous fashion. In order to evaluate `hobf` we then simply provide the name of a binary boolean function and two boolean values as

```
1 >hobf xor 0 0 = 0
  >hobf xor 0 1 = 1
3 >hobf xor 1 0 = 1
  >hobf xor 1 1 = 0
```

Listing 3: Example output of `hobf`

Assessment of abstraction potential in NLDs

An important consideration is to state precisely some criterion for determining that a particular NLD has achieved potential for an explicit form of abstraction such as is indicated in the relationship between expressions 3 and 4. In abstracting the binary boolean operator \wedge to the functional variable f , which is the fundamental transformation enabling the derivation of 4 from 3, we imply that any physical implementation must take at least three rather than two inputs and the first of these must specify a particular binary boolean operator to apply to the latter two boolean input values. This roughly means that any system that at least partially implements the lambda expression or function specified in equation 4 and listing 1 respectively, must be capable of interpreting the concept of “selection from a set” whose size is determined by the subset of the 16 binary boolean operators that is already implemented in a lower-level form. Indeed another representation of equation 4 as a partial or total set function could be written as $hobf : Hex \times Bool \times Bool \rightarrow Bool$ where we interpret $Bool$ and Hex as two and sixteen element sets respectively. This point of view makes clear that the capacity for selection from two element sets is already apparent in the binary boolean operator written as a set function $\wedge : Bool \rightarrow Bool$. The difference between these is that the system must implement the typing constraints necessary to distinguish a set that takes on sixteen possible values from one that takes on two. For example, in the application of the function $hobf$ the following evaluate as expected given the definition: $hobf \wedge 1 0 = 0$ or $hobf \wedge 1 1 = 1$. However, what is to be expected given inputs such as: $hobf 1 \wedge 0$ or $hobf 1 1 \wedge$? In these cases an output type intuitively associated to *error* is required to indicate that the realization of a system implementing equation

4 has indeed correctly implemented the necessary typing constraints to claim that the function has been realized. In the case of neuronal logic devices, this alternative output should differ in a measurable way from those associated to 1 and 0.

Utilizing synaptic plasticity to generate higher order NLDs

In order to take advantage of the immense capacity for computation available in biological systems and neural networks, the biological complexity of each neuron and synaptic connection must be taken into consideration. Namely, the presence of paired pulse depression and facilitation, which has been previously demonstrated in cultured hippocampal neurons [6], could be exploited to create neural objects with similar design, but more complex behavior than existing NLDs. Specifically, if paired pulse coordination can be shown to exist, especially in the forward direction more significantly than the reverse, then perhaps chains of action potentials could be used as the logical readouts instead of single action potentials. If this is the case, then something approaching a hexadecimal computing system can be created using these logical elements, which would allow for significant information compression. In the long term, something approximating the design of a “neural VLSI” supporting a full logic system such as the simply typed lambda calculus could be implemented.

The experiments described in the previous section allow for the engineering of a rich and complex variety of logical device configurations. The vast richness that the theoretical considerations will enable ensures a steady flow of experimental systems that can be designed quickly to test the theoretical ideas. In that sense, the cultured living neuronal networks supply a never-ending ‘Lego’ toolbox for verifying and evaluating different theoretical strategies for engineered networks. The numerous parameters that can be varied and checked allow the different conceptual approaches to biological computation to be tested.

From NLDs to principles of cognition

From the point of view that identifies the capacity for abstraction with computation, the investigation of neuronal logic devices capable of such function provides a framework in which various hypotheses from cognitive science could begin to be evaluated at the level of well-defined neural circuits. By attempting to isolate minimal implementation criteria, this approach may serve to complement, enable simpler explanations of, or identify paradoxical results derived from studies that treat whole brains and their associated sensory apparatus as their object of study [7,8].

As noted above, we associate the failure of a cultured neuronal network to produce a valid computation with the fact that it is grown out of context, out of its natural surroundings. By geometric constraints, we have shown that we can coax some of the networks connections to go in the direction we engineer, and a modicum of computation is restored. However, the idea that the ‘software’ is growing without the presence of the ‘hardware’ it is accustomed to, leads us to the idea of co-culturing a different set of neurons along with the regular hippocampal ones. The obvious choice for us is that of sensory neurons, since those are the subset of neurons that communicate between the body of the organism and its brain. In this way we hope to recreate an information channel that exists in the developing brain, and allow the neurons in the animal to attain functional connectivity according to the inputs it receives. We are thus communicating with a number of groups that study pain, with the intention of extracting neurons that can synapse onto the CNS one and convey the signals associated with stress and pain. This seems a reasonable initial choice, since the signal of pain is one that a neuron should be designed not to ignore.

Section 4: Role and expertise of the PIs

Prof. Aviv Bergman will be responsible for the theoretical arm of this project. Together with 3 of his Ph.D. Students, Cameron Smith, Daniel Biro and Ximo Pechaun, they will develop the computational platform and analytical tools to predict neuronal architecture which in turn will help guide the construction of environmental conditions to address desirable computational tasks to be performed by in vitro neuronal networks. Prof. Bergman's expertise spans a number of theoretical areas including Artificial Neural Networks, Dynamical Systems, Mathematical Evolutionary Biology, and Theoretical Systems Biology.

Prof. Elisha Moses ...

Section 5: Educational involvement

In the Bergman lab, three students will be involved in developing the theoretical aspect of this project in close interaction with the PI. These three students have a complementary set of backgrounds that will be crucial to take advantage of in the development of the research program outlined here. Two of the students have broad theoretical background in mathematics, computer science, engineering, and have each completed the standard courses for both the Department of Systems and Computational Biology Ph.D. program and for the M.D. program at the Albert Einstein College of Medicine. A third student has a combined theoretical and experimental background in chemistry, physics and biology. These three students share an office and all have benches together in the Bergman wet lab. We will take full advantage of this broad and deep knowledge base.

In the Moses lab ...

Students and postdoctoral fellows in the Bergman and Moses labs will interact on a weekly basis via video conference. More extensive interaction will be fostered by an exchange program that we plan to engage in at important points in the project for the purpose of increasing the level of coherence between the theoretical and experimental aspects of this work. This exchange program will expose the experimentalists in the Moses lab to the theoretical and experimental work in the Bergman lab and expose theoreticians in the Bergman lab to the combined theoretical and experimental work in the Moses lab. Despite the fact that both labs already combine theory and experiment, they do so in very different ways and exposure to each of these models for combining theory and experiment will be crucial for these students as they move on to make very difficult decisions in their careers with respect to postdoctoral training and starting their own labs.

References Cited

- [1] J. B. Tenenbaum, C. Kemp, T. L. Griffiths, and N. D. Goodman. How to Grow a Mind: Statistics, Structure, and Abstraction. *Science*, 331(6022):1279–1285, March 2011.
- [2] Ofer Feinerman, Assaf Rotem, and Elisha Moses. Reliable neuronal logic devices from patterned hippocampal cultures. *Nature Physics*, 4(12):967–973, October 2008.
- [3] Stephen A. Ward and Robert H. Halstead. *Computation Structures*. The MIT Press, 1989.
- [4] Harold Abelson and Gerald Jay Sussman. *Structure and Interpretation of Computer Programs - 2nd Edition (MIT Electrical Engineering and Computer Science)*. The MIT Press, 1996.
- [5] H.P. Barendregt. *The Lambda Calculus, Its Syntax and Semantics*. North Holland, 1985.
- [6] D D Cummings, K S Wilcox, and M a Dichter. Calcium-dependent paired-pulse facilitation of miniature EPSC frequency accompanies depression of EPSCs at hippocampal synapses in culture. *The Journal of neuroscience*, 16(17):5312–23, September 1996.
- [7] James L McClelland, Matthew M Botvinick, David C Noelle, David C Plaut, Timothy T Rogers, Mark S Seidenberg, and Linda B Smith. Letting structure emerge: connectionist and dynamical systems approaches to cognition. *Trends in cognitive sciences*, 14(8):348–56, August 2010.
- [8] Thomas L Griffiths, Nick Chater, Charles Kemp, Amy Perfors, and Joshua B Tenenbaum. Probabilistic models of cognition: exploring representations and inductive biases. *Trends in cognitive sciences*, 14(8):357–64, August 2010.

Biographical Sketch: Your Name